

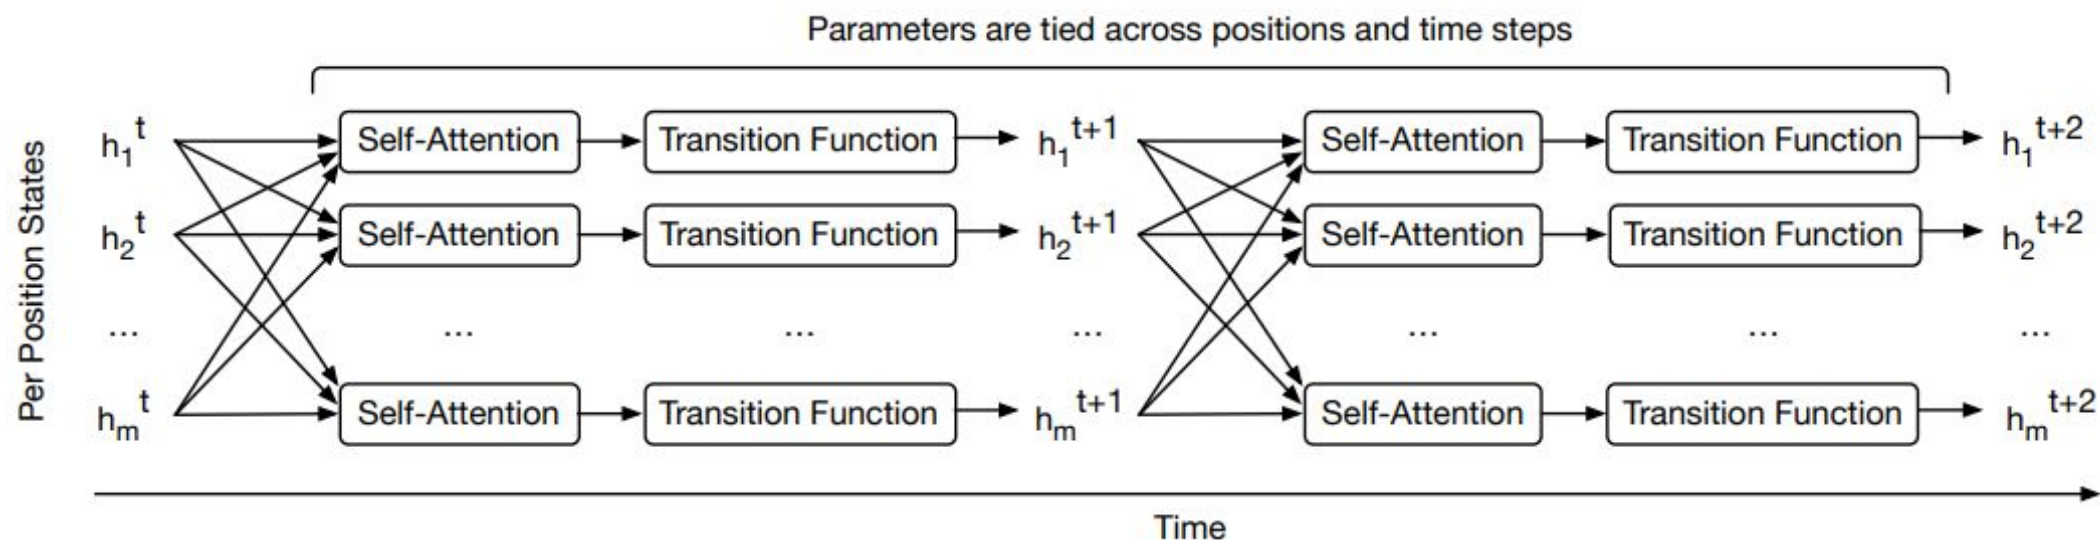
Universal Transformers

presenter: Baosong Yang

Motivation

- Failing to generalize in many tasks
- Not computationally universal, limiting its theoretical expressivity
- Core difference to Transformer:
 - Repeatedly revises its representations of all symbols in the sequence with each recurrent step. (Using RNN to replace the stacked architecture)
 - Adopting ACT (adaptive computation time)

Model



- Recurrent transition function
 - either a separable convolution
 - or a fully-connected neural network that consists of a single rectified-linear activation function between two affine transformations
- Parameters are tied across time steps (layers)

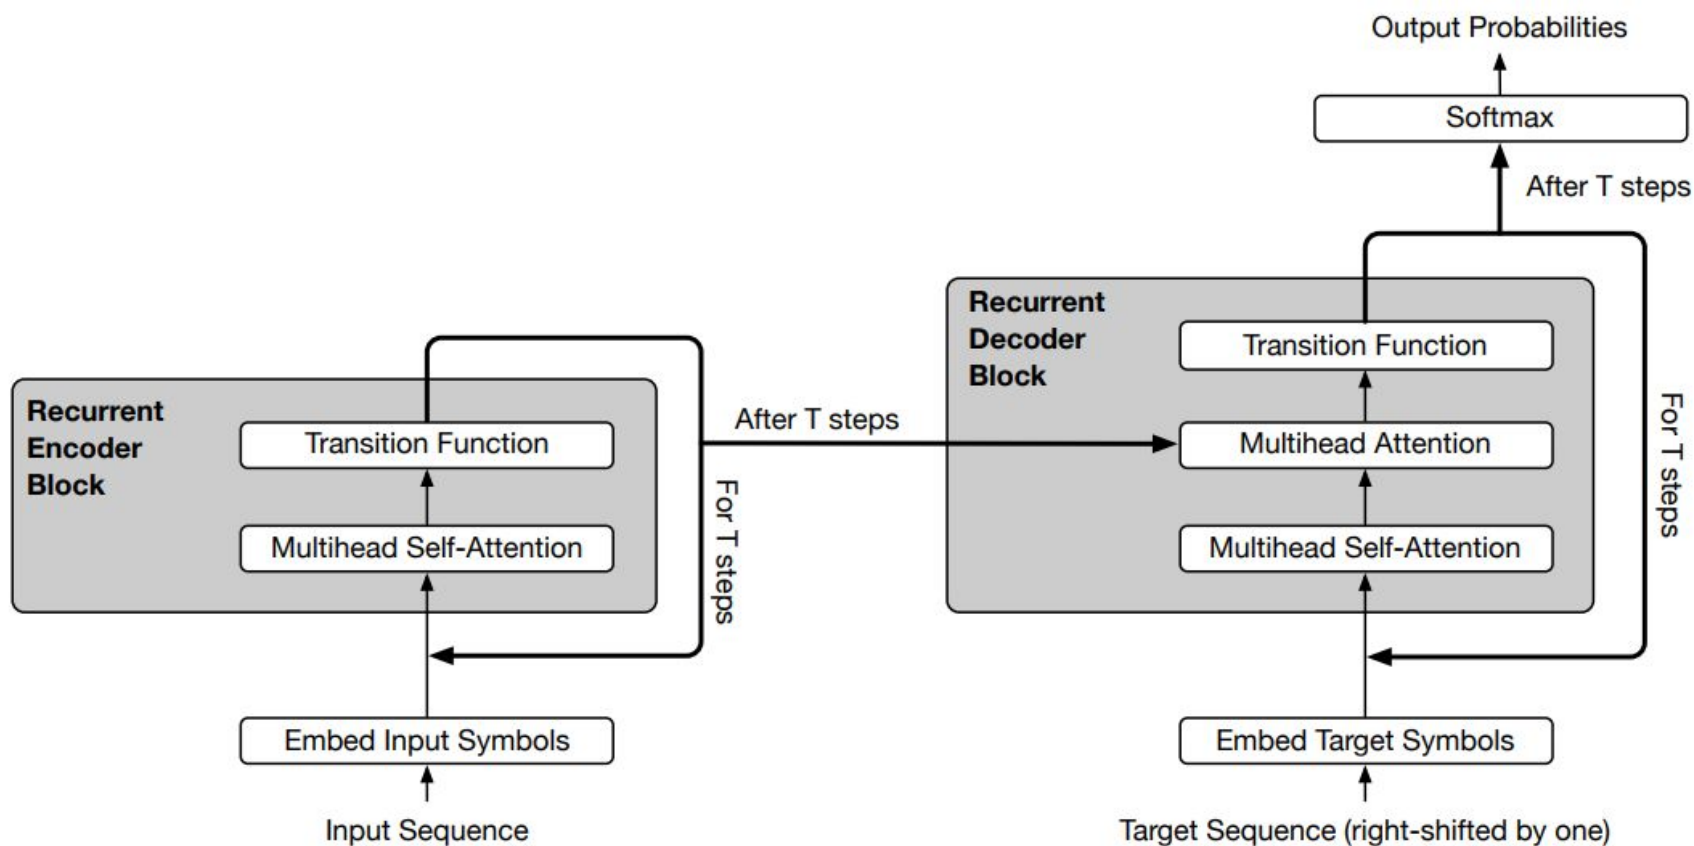
Model

$$H^t = \text{LayerNorm}(A^{t-1} + \text{Transition}(A^t))$$

$$\text{where } A^t = \text{LayerNorm}(H^{t-1} + \text{MultiHeadSelfAttention}(H^{t-1} + P^t)),$$

$$P_{pos,2j}^t = \sin(pos/10000^{2j/d}) \oplus \sin(t/10000^{2j/d})$$

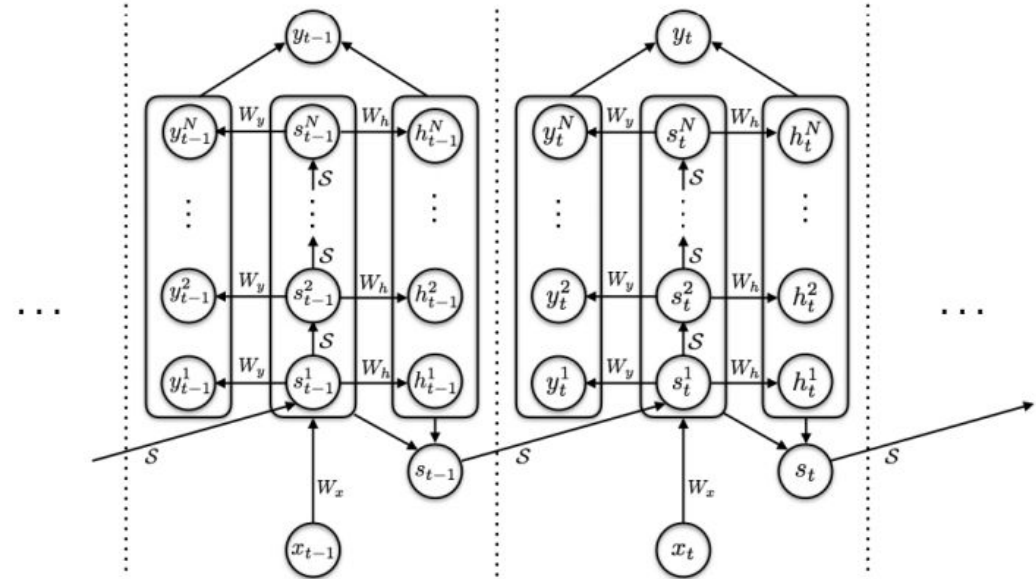

$$P_{pos,2j+1}^t = \cos(pos/10000^{2j/d}) \oplus \cos(t/10000^{2j/d}).$$



Adaptive Computation Time

- Dynamically modulating the number of computational steps needed to process each input symbol.
- Based on a scalar pondering value predicted by the model at each step.
- Once the per-symbol recurrent block halts, its state is simply copied to the next step until all blocks halt, or we reach a maximum number of steps.

Adaptive Computation Time



- input: the hidden state of the t-th element in n-th layer
- output: probability of halting

$$h_t^n = \sigma(W_h s_t^n + b_h) \quad p_t^n = \begin{cases} R(t) & \text{if } n = N(t) \\ h_t^n & \text{otherwise} \end{cases} \quad N(t) = \min\{n' : \sum_{n=1}^{n'} h_t^n \geq 1 - \epsilon\}, \quad R(t) = 1 - \sum_{n=1}^{N(t)-1} h_t^n$$

- Differentiable: $s_t = \sum_{n=1}^{N(t)} p_t^n s_t^n \quad y_t = \sum_{n=1}^{N(t)} p_t^n y_t^n$

Experiments

- bABI QA: The dataset consists of 20 different tasks. Given a story, answer questions.

Model	10K examples		1K examples	
	train single	train joint	train single	train joint
Previous best results:				
QRNet [24]	0.3 (0/20)	-	-	-
Sparse DNC [23]	-	2.9 (1/20)	-	-
GA+MAGE [8]	-	-	8.7 (5/20)	-
MemN2N [26]	-	-	-	12.4 (11/20)
Our Results:				
Transformer [31]	15.2 (10/20)	22.1 (12/20)	21.8 (5/20)	26.8 (14/20)
Universal Transformer (this work)	0.23 (0/20)	0.47 (0/20)	5.31 (5/20)	8.50 (8/20)
Adapt. Univ. Transformer (this work)	0.21 (0/20)	0.29 (0/20)	4.56 (3/20)	7.85 (5/20)

Table 1: Average error and number of failed tasks ($> 5\%$ error) out of 20 (in parentheses; lower is better in both cases) on the bAbI dataset under the different training/evaluation setups. We indicate state-of-the-art where available for each, or ‘-’ otherwise.

Experiments

- Subject-Verb Agreement: Number agreement between subject and verb

Model	Number of attractors						Total
	0	1	2	3	4	5	
Previous best results [33]:							
Best Stack-RNN	0.994	0.979	0.965	0.935	0.916	0.880	0.9923
Best LSTM	0.993	0.972	0.95	0.922	0.900	0.842	0.9911
Best Attention	0.994	0.977	0.959	0.929	0.907	0.842	0.9918
Our results:							
Transformer	0.9733	0.9412	0.9316	0.9167	0.9014	0.8834	0.9616
Universal Transformer	0.9934	0.9712	0.9690	0.9400	0.9206	0.8915	0.9917
Adapt. Univ. Transf. (small)	0.9932	0.9801	0.9714	0.9608	0.9521	0.9314	0.9920
Adapt. Univ. Transf. (base)	0.9943	0.9720	0.9516	0.9567	0.9314	0.9034	0.9924

Table 2: Accuracy on the subject-verb agreement number prediction task (higher is better).

Experiments

- The LAMBADA task is a language modeling task consisting of predicting a missing target word given its (target) sentence and a broader context of 4-5 preceding sentences.

Model	LM Perplexity & (Accuracy)			RC Accuracy		
	control	dev	test	control	dev	test
Neural Cache [10]	129	139	-	-	-	-
Dhingra et al. [7]	-	-	-	-	-	0.5569
Transformer	154 (0.14)	5336 (0.0)	9725 (0.0)	0.4102	0.4401	0.3988
LSTM	138 (0.23)	4966 (0.0)	5174 (0.0)	0.1103	0.2316	0.2007
Universal Transformer	131(0.32)	279 (0.18)	319 (0.17)	0.4801	0.5422	0.5216
Adaptive Universal Transformer	130 (0.32)	135 (0.22)	142 (0.19)	0.4603	0.5831	0.5625

Table 3: LAMBADA language modeling (LM) perplexity (lower better) with accuracy in parentheses (higher better), and Reading Comprehension (RC) accuracy results (higher better). ‘-’ indicates no reported results in that setting.

Experiments

- Algorithmic Tasks:

Model	Copy		Reverse		Addition	
	char-acc	seq-acc	char-acc	seq-acc	char-acc	seq-acc
LSTM	0.45	0.09	0.66	0.11	0.08	0.0
Transformer	0.53	0.03	0.13	0.06	0.07	0.0
Universal Transformer	0.91	0.35	0.96	0.46	0.34	0.02
Neural GPU*	1.0	1.0	1.0	1.0	1.0	1.0

Table 4: Accuracy (higher better) on the algorithmic tasks, trained on decimal strings of length 40 and evaluated on length 400 from [17]. *Note that the Neural GPU was trained with a special curriculum to obtain the perfect result, while other models are trained without any curriculum.

Experiments

- Learning to Execute:

Model	Copy		Double		Reverse	
	char-acc	seq-acc	char-acc	seq-acc	char-acc	seq-acc
LSTM	0.78	0.11	0.51	0.047	0.91	0.32
Transformer	0.98	0.63	0.94	0.55	0.81	0.26
Universal Transformer	1.0	1.0	1.0	1.0	1.0	1.0

Table 5: Character-level (*char-acc*) and sequence-level accuracy (*seq-acc*) results on the Memorization LTE tasks, with maximum length of 55.

Model	Program		Control		Addition	
	char-acc	seq-acc	char-acc	seq-acc	char-acc	seq-acc
LSTM	0.53	0.12	0.68	0.21	0.83	0.11
Transformer	0.71	0.29	0.93	0.66	1.0	1.0
Universal Transformer	0.89	0.63	1.0	1.0	1.0	1.0

Table 6: Character-level (*char-acc*) and sequence-level accuracy (*seq-acc*) results on the Program Evaluation LTE tasks with maximum nesting of 2 and length of 5.

Experiments

- Machine Translation:

Model	BLEU
Universal Transformer <i>small</i>	26.8
Transformer <i>base</i> [31]	28.0
Weighted Transformer <i>base</i> [1]	28.4
Universal Transformer <i>base</i>	28.9

Conclusion

- An new architecture for Transformer.
- More parameters or more biases?
- Theoretical support for context-specific?