

# Paper Reading: Neural Machine Translation with Gumbel-Greedy Decoding (AAAI 2018)

Jiatao Gu, Daniel Jiwoong Im, Victor O.K. Li

The University of Hong Kong

# Motivation


- ▶ Once the model is trained, the most probable output which maximum the log-likelihood during training cannot be properly found at the test time.
- ▶ Avoid solving the maximum a posteriori problem over translation sentences at test phase.

# Discriminator-Generator framework

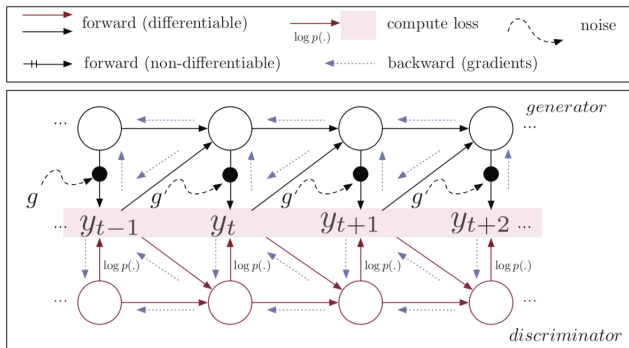
- ▶ NMT-discriminator: measures the log-likelihood at word level  $-\log_{p_\theta}(Y | X)$  given the source sentence  $X$  and a translation  $Y$ .
- ▶ NMT-generator: generates the translation by taking the output of the word as an input to next step recursively,  $Y = G_\theta(X)$ , given the source sentence  $X$ . ( $G$  is usually a search-based method).
- ▶ NMT discriminator's score:

$$J(\phi) = \mathbb{E}_{Y \sim G_\phi} \log p_\theta(Y|X)$$

# Discriminator, Generator and Gumbel-Greedy Decoding



(a)



(b)

# Gumbel-Greedy Decoding

- ▶ The Gumbel-Max Trick:

$$y = \operatorname{argmax} (g + a), g \sim \text{Gumbel i.i.d.}$$

Where each element in  $g$  can be computed using the inverse transform sampling of an auxiliary random uniform variable  $u_i \sim U(0, 1)$ ,  $g_i = -\log(-\log(u_i))$

- ▶ Gumbel-Softmax Relaxation:

$$\hat{y} = \operatorname{softmax}((g + a)/\tau), g \sim \text{Gumbel i.i.d.}$$

Where  $\tau \in (0, \infty)$ . The softmax function approaches argmax operations as  $\tau \rightarrow 0$ , and it becomes uniform when  $\tau \rightarrow \infty$ .

- ▶ Stright-Throught(ST) Gumbel: During the forward phase, use the Gumbel-max, while computing the gradient of the Gumbel-softmax.

► Arbitrary Decoding Algorithms as Gumbel-Greedy Decoding:

$$y = \operatorname{argmax} (g + a), g \sim Q$$

$$g_i^* = \begin{cases} g', & y_i \text{ is selected} \\ g' - \log \left[ 1 + e^{g' - \tilde{g}_i} \right], & \text{otherwise} \end{cases} \quad (12)$$

where

$$\tilde{g}_i = -\log(-\log(u_i)) + a_i, u_i \sim \mathcal{U}(0, 1)$$

and the “top-gumbel”

$$g' = -\log(-\log(u)) + \log \left( \sum_i \exp(a_i) \right), u \sim \mathcal{U}(0, 1)$$

# Gumbel-Greedy Decoding Algorithm

## ► With Regularization:

$$\mathbb{E}_{G_\phi} [\log p_\theta(Y|X)] - \mathbb{E}_{G_\phi} [\log p_{\phi'}(Y|X)] \quad (13)$$

where we use  $\phi'$  to represent a copy of the current parameters  $\phi$  and make it as a “discriminator”. Note that gradients w.r.t  $\phi$  will not flow into  $\phi'$ .

## ► Adversarial Learning:

$$\mathbb{E}_D [\log p_\theta(Y|X)] - \mathbb{E}_G [\log p_\theta(Y|X)] \quad (14)$$

where  $D$  is the empirical distribution of real translation. In practice, we alternate the training of the generator and the discriminator iteratively.

---

### Algorithm 1 Gumbel-Greedy Decoding

---

**Require:** discriminator  $p_\theta$ , generator  $G_\phi$ ,  $N_d \geq 0$ ,  $N_g > 0$


- 1: Train  $\theta$  using MLE/REINFORCE on training set  $D$ ;
- 2: Initialize  $\phi$  using  $\theta$ ;
- 3: Shuffle  $D$  twice into  $D_\theta$  and  $D_\phi$
- 4: **while** stopping criterion is not met **do**
- 5:   **for**  $t = 1 : N_g$  **do**     *// learn the generator*
- 6:     Draw a translation pair:  $(X, \_) \sim D_\phi$ ;
- 7:     Obtain  $Y, \hat{Y} = \text{GUMBELDEC}(G, X)$
- 8:     Compute forward pass  $\sim X, Y$  with Eq. 13
- 9:     Compute backward pass  $\sim X, \hat{Y}$ , update  $\phi$
- 10:   **for**  $t = 1 : N_d$  **do**     *// learn the discriminator*
- 11:     Draw a translation pair:  $(X, Y^*) \sim D_\theta$ ;
- 12:     Obtain  $Y, \_ = \text{GUMBELDEC}(G, X)$
- 13:     Compute forward pass  $\sim X, Y, Y^*$  with Eq. 14
- 14:     Compute backward pass  $\sim X, Y, Y^*$ , update  $\theta$

**Function:** GUMBELDEC( $G, X$ )


- 1: **if**  $G = \text{'sampling'}$  **then**
  - 2:   Sample  $g \sim \text{Gumbel i.i.d.}$
  - 3:   Obtain  $Y, \hat{Y}$  with Eq. 7 and Eq. 8
  - 4: **else**
  - 5:   Obtain  $Y = G(X)$
  - 6:   Infer  $g$  with Eq. 12
  - 7:   Obtain  $Y, \hat{Y}$  with Eq. 8
  - 8: **Return**  $Y, \hat{Y}$
-




# Experiments




(a) comparison with variant temperature



(b) comparison w/o entropy regularization



(c) comparison of greedy decoding



(d) comparison of sampling

	Model	DE-EN	EN-DE	CS-EN	EN-CS
Greedy	MLE	21.63	18.97	18.90	14.49
	RL	22.56	19.32	19.45	15.02
	GGD	<b>23.27</b>	<b>19.81</b>	<b>20.62</b>	<b>16.04</b>
Beam	MLE	24.46	21.33	21.20	16.20
	RL	25.12	<b>22.13</b>	21.92	17.02
	GGD	<b>25.32</b>	21.97	<b>22.47</b>	<b>17.64</b>

Table 1: The greedy decoding and the beam-search performance of models trained with GGD-GAN against MLE and REINFORCE (referred to RL). BLEU scores are calculated on the test sets.

# Conclusion

- ▶ Use the Gumbel-Softmax reparameterization trick to make the generative network differentiable and can be trained through standard stochastic gradient methods.