

Approaches for Subword Regularization

Yongchang Hao

Aug 6, 2020

Neural Machine Translation of Rare Words with Subword Units

A quick review

- The standard subword model: BPE
- At each merge steps:
 1. count current pairs in dict
 2. merge the most frequent pair in dict
- BPE is deterministic

Algorithm 1 Learn BPE operations

```
import re, collections

def get_stats(vocab):
    pairs = collections.defaultdict(int)
    for word, freq in vocab.items():
        symbols = word.split()
        for i in range(len(symbols)-1):
            pairs[symbols[i], symbols[i+1]] += freq
    return pairs

def merge_vocab(pair, v_in):
    v_out = {}
    bigram = re.escape(' '.join(pair))
    p = re.compile(r'(?!\S)' + bigram + r'(?!\S)')
    for word in v_in:
        w_out = p.sub(' '.join(pair), word)
        v_out[w_out] = v_in[word]
    return v_out

vocab = {'l o w </w>' : 5, 'l o w e r </w>' : 2,
         'n e w e s t </w>':6, 'w i d e s t </w>':3}
num_merges = 10
for i in range(num_merges):
    pairs = get_stats(vocab)
    best = max(pairs, key=pairs.get)
    vocab = merge_vocab(best, vocab)
    print(best)
```

Subword Regularization: Improving Neural Network Translation Models with Multiple Subword Candidates - Taku Kudo - ACL 2018

Intuition

- Use the segmentation ambiguity as noise to improve robustness

- | Subwords (- means spaces) | Vocabulary id sequence |
|---------------------------|-------------------------|
| _Hell/o/_world | 13586 137 255 |
| _H/ello/_world | 320 7363 255 |
| _He/llo/_world | 579 10115 255 |
| _/He/l/l/o/_world | 7 18085 356 356 137 255 |
| _H/el/l/o/_/world | 320 585 356 137 7 12295 |

Table 1: Multiple subword sequences encoding the same sentence “Hello World”

Subword Regularization: Improving Neural Network Translation Models with Multiple Subword Candidates

Approach - Training

- Induce the segmentation distribution P_d to the original loss

- $$\mathcal{L}_m(\theta) = \sum_{s=1}^{|D|} \mathbb{E}_{\substack{x \sim P_d(x | X^{(s)}) \\ y \sim P_d(y | Y^{(s)})}} [\log P(y | x; \theta)]$$

- Connect the loss function with subword segmentation

Subword Regularization: Improving Neural Network Translation Models with Multiple Subword Candidates

Approach - Inference

- one-best decoding:
 - Only translate $\operatorname{argmax}_x P_d(x | X)$
- n-best decoding:
 - Translate n best $P(x | X)$, select one with maximal $\frac{\log P(y | x)}{|y|^\lambda}$

Subword Regularization: Improving Neural Network Translation Models with Multiple Subword Candidates

Main result

Corpus	Language pair	baseline (BPE)	Proposed (one-best decoding)			Proposed (n -best decoding, $n=64$)		
			$l=1$	$l=64$ $\alpha=0.1$	$l=\infty$ $\alpha=0.2/0.5$	$l=1$	$l=64$ $\alpha=0.1$	$l=\infty$ $\alpha=0.2/0.5$
IWSLT15	en \rightarrow vi	25.61	25.49	27.68*	27.71*	25.33	28.18*	28.48*
	vi \rightarrow en	22.48	22.32	24.73*	26.15*	22.04	24.66*	26.31*
	en \rightarrow zh	16.70	16.90	19.36*	20.33*	16.73	20.14*	21.30*
	zh \rightarrow en	15.76	15.88	17.79*	16.95*	16.23	17.75*	17.29*
IWSLT17	en \rightarrow fr	35.53	35.39	36.70*	36.36*	35.16	37.60*	37.01*
	fr \rightarrow en	33.81	33.74	35.57*	35.54*	33.69	36.07*	36.06*
	en \rightarrow ar	13.01	13.04	14.92*	15.55*	12.29	14.90*	15.36*
	ar \rightarrow en	25.98	27.09*	28.47*	29.22*	27.08*	29.05*	29.29*
KFTT	en \rightarrow ja	27.85	28.92*	30.37*	30.01*	28.55*	31.46*	31.43*
	ja \rightarrow en	21.37	21.46	22.33*	22.04*	21.37	22.47*	22.64*
ASPEC	en \rightarrow ja	40.62	40.66	41.24*	41.23*	40.86	41.55*	41.87*
	ja \rightarrow en	26.51	26.76	27.08*	27.14*	27.49*	27.75*	27.89*
WMT14	en \rightarrow de	24.53	24.50	25.04*	24.74	22.73	25.00*	24.57
	de \rightarrow en	28.01	28.65*	28.83*	29.39*	28.24	29.13*	29.97*
	en \rightarrow cs	25.25	25.54	25.41	25.26	24.88	25.49	25.38
	cs \rightarrow en	28.78	28.84	29.64*	29.41*	25.77	29.23*	29.15*

Adversarial Subword Regularization for Robust Neural Machine Translation - Jungsoo Park, Mujeen Sung, Jinhyuk Lee, Jaewoo Kang - arXiv 2020

Intuition

- Generate subword segmentation that maximize the loss
- $\hat{x} = \operatorname{argmax}_{x \in \Omega(X_i)} \left((\nabla_{\tilde{x}} \mathcal{L})^T \cdot (emb(x) - emb(\tilde{x})) \right)$
- Ω denotes all possible segmentation, x denotes a real word where consists of subwords $x = [x_1, \dots, x_k]$, $emb(x) = \frac{1}{k} \sum_{i=1}^k emb(x_i)$
- $x \in \Omega(X_i)$ is done by Kudo's work.

Adversarial Subword Regularization for Robust Neural Machine Translation

Main result

Lang Pair	BASE	SR	ADVSR
IWSLT17			
FR → EN	37.9	38.1	38.7
EN → FR	38.8	39.1	39.9
AR → EN	31.7	32.3	33.3
EN → AR	14.4	14.3	14.7
IWSLT15			
CS → EN	28.9	30.5	32.0
EN → CS	20.4	21.7	23.9
VI → EN	28.1	28.4	29.0
EN → VI	30.9	31.7	32.3
IWSLT13			
PL → EN	19.1	19.7	20.9
EN → PL	13.5	14.1	15.1
TR → EN	21.3	22.6	23.4
EN → TR	12.6	14.4	14.0

Adversarial Subword Regularization for Robust Neural Machine Translation

Robustness analysis

Dataset	BASE	SR	ADVSR
MTNT2018			
FR \rightarrow EN	25.7	27.6	27.2
EN \rightarrow FR	26.7	27.5	28.2
MTNT2018 + FT			
FR \rightarrow EN	36.5	37.9	38.8
EN \rightarrow FR	33.2	34.4	35.3
MTNT2019			
FR \rightarrow EN	27.6	29.3	30.2
EN \rightarrow FR	22.8	23.8	24.1
MTNT2019 + FT			
FR \rightarrow EN	36.2	38.1	38.6
EN \rightarrow FR	27.6	28.2	28.9

Table 3: BLEU scores on the out-of-domain MTNT dataset. **FT** denotes finetuning with the MTNT2018 training dataset.

Method	0.1	0.2	0.3	0.4	0.5
FR \rightarrow EN					
BASE	30.7	25.6	20.3	16.2	11.4
SR	33.2	28.5	23.3	18.7	14.7
ADVSR	34.8	32.0	29.2	25.7	22.2
EN \rightarrow FR					
BASE	31.1	24.2	18.6	14.6	10.6
SR	34.2	27.8	23.9	18.9	14.4
ADVSR	35.1	30.3	26.4	23.0	19.1

Table 4: BLEU scores on the synthetic dataset of typos. The column lists results for different noise fractions.

Drawbacks of these 2 papers

Too complicated

- EM algorithm to optimize vocabulary
- Viterbi algorithm to find best segmentation
- Enhanced Suffix Array to find frequent substring
- Forward-DP Backward-A* in l-best search
- Forward-Filtering Backward-Sampling for infinity search
-

BPE-Dropout: Simple and Effective Subword Regularization - Ivan Provilkov, Dmitrii Emelianenko, Elena Voita - ACL 2020

Approach

- At each merge steps:
 1. count current pairs in dict
 2. randomly drop some merge operations with the probability p
 3. merge the most frequent pair in dict
- BPE-dropout is stochastic

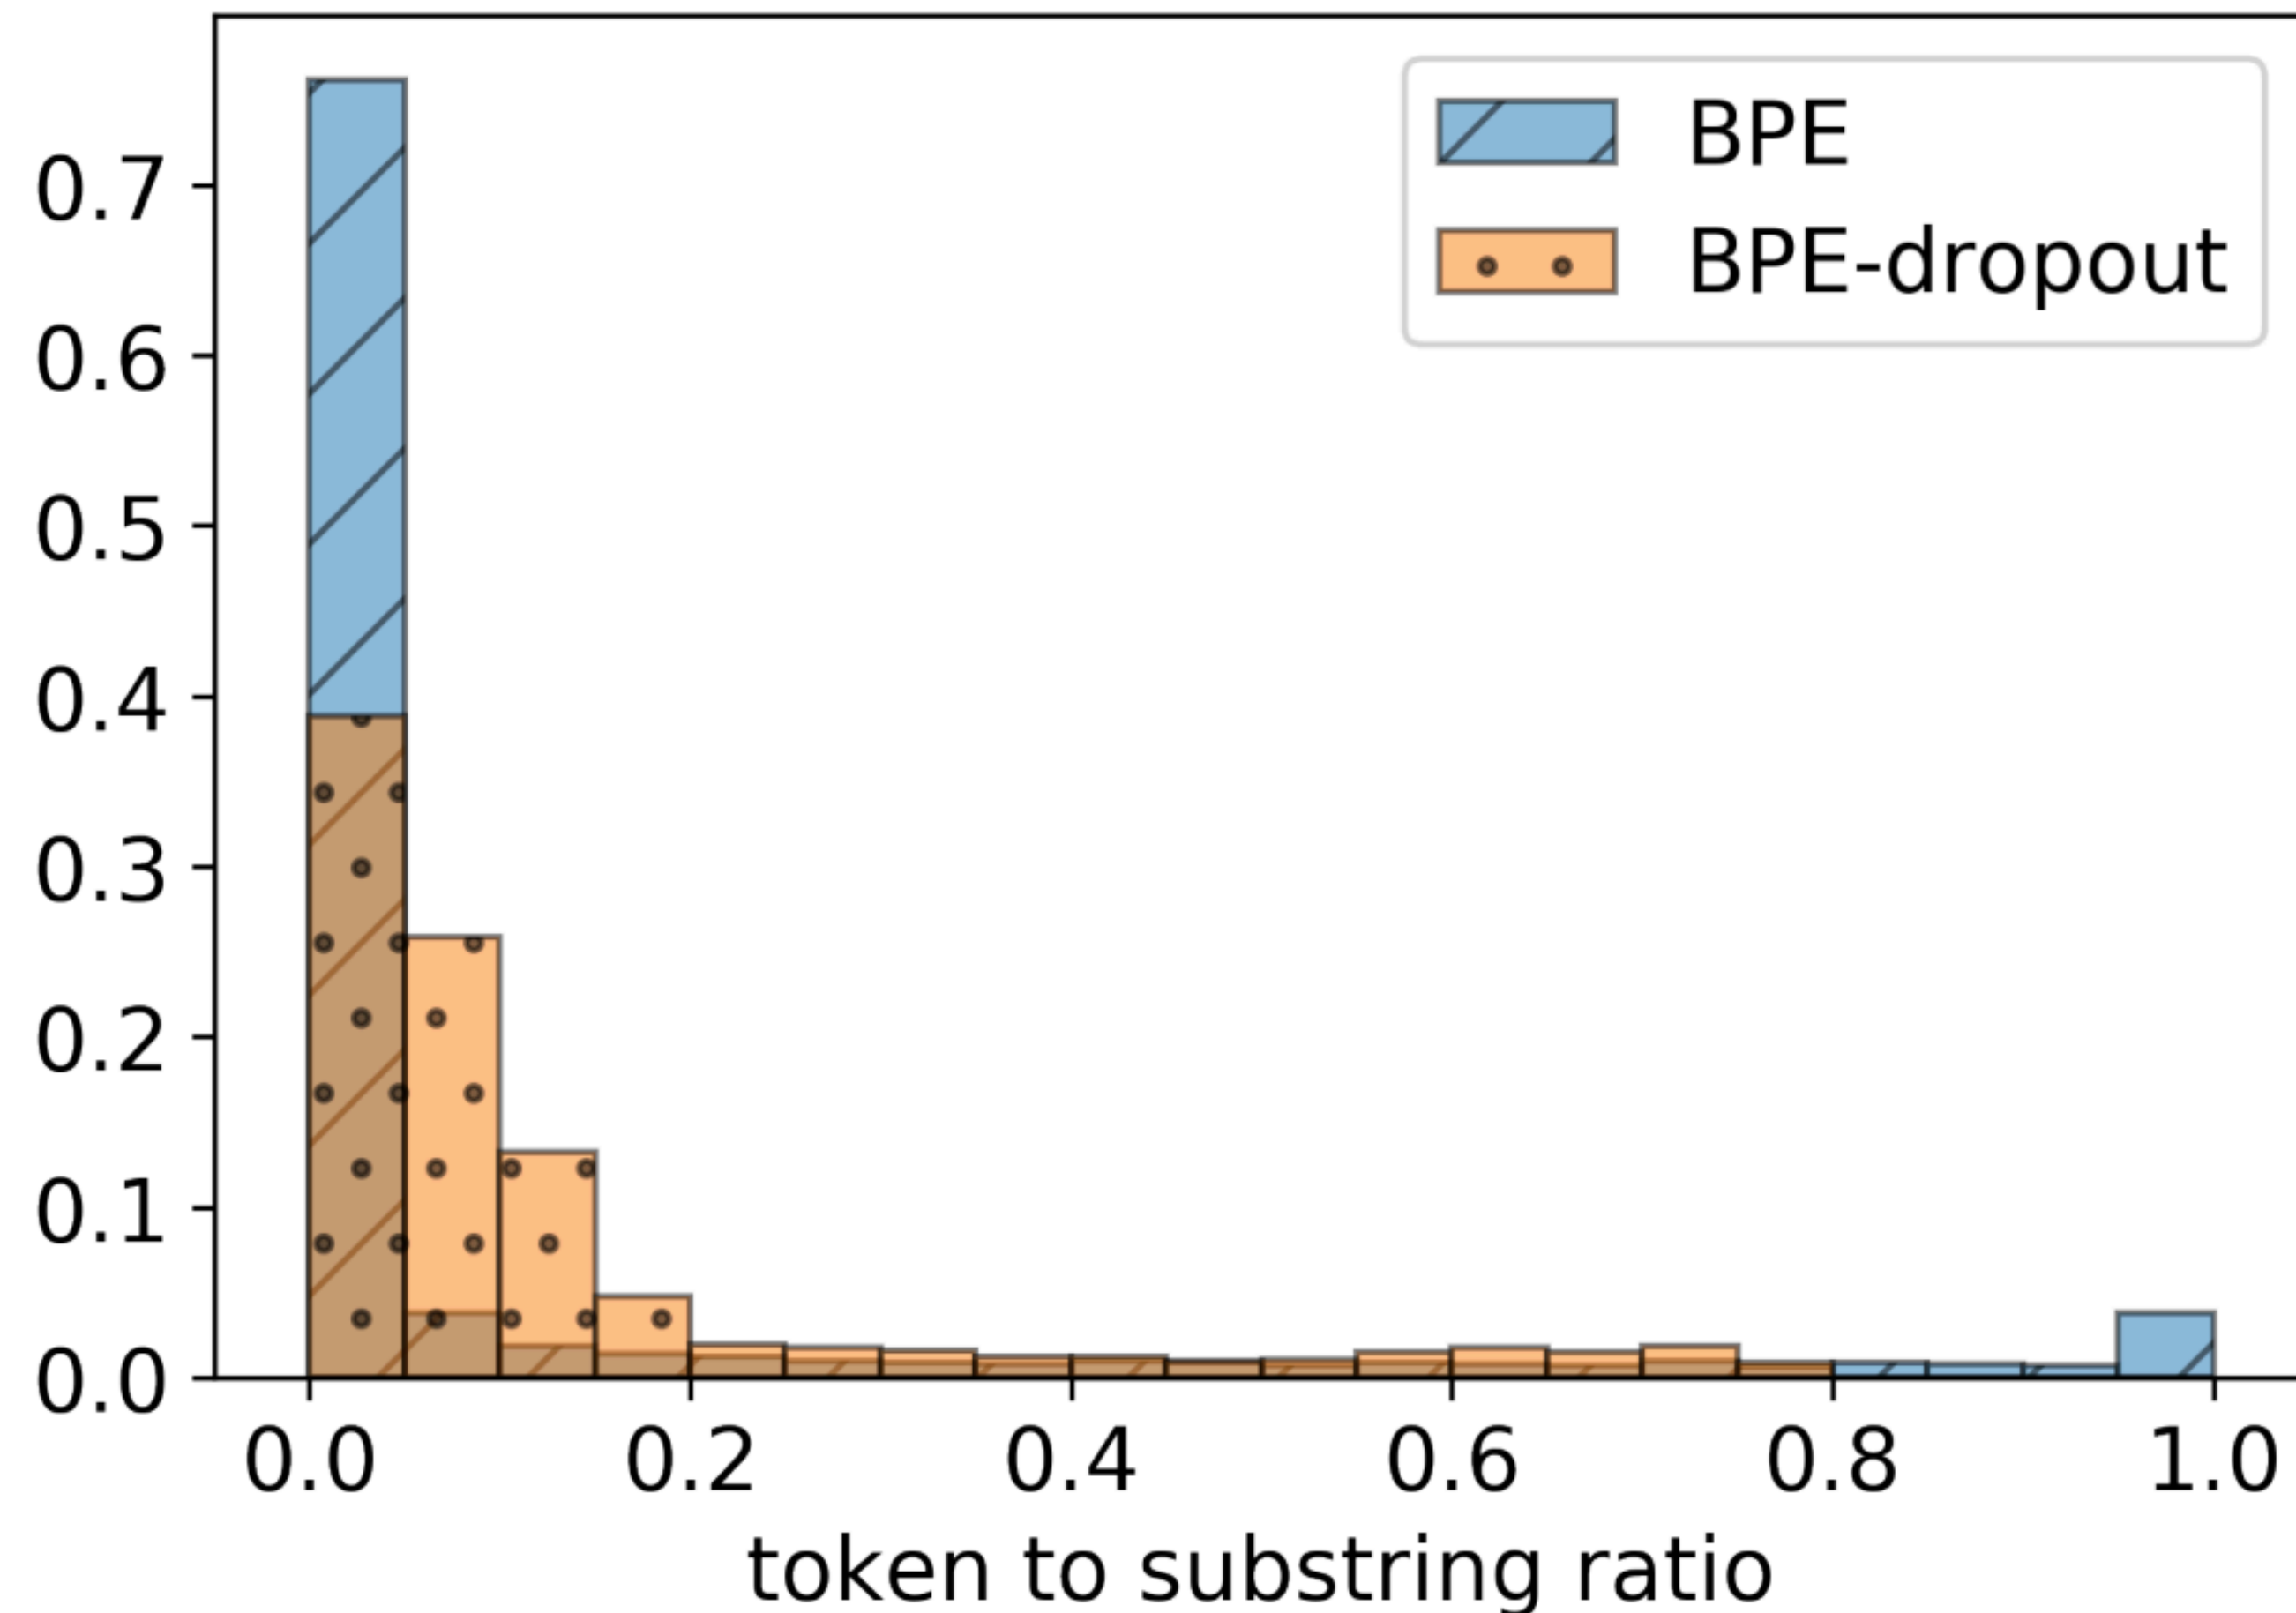
BPE-Dropout: Simple and Effective Subword Regularization

Main result

	BPE	Kudo (2018)	<i>BPE-dropout</i>
IWSLT15			
En-Vi	31.78	32.43	33.27
Vi-En	30.83	32.36	32.99
En-Zh	20.48	23.01	22.84
Zh-En	19.72	21.10	21.45
IWSLT17			
En-Fr	39.37	39.45	40.02
Fr-En	38.18	38.88	39.39
En-Ar	13.89	14.43	15.05
Ar-En	31.90	32.80	33.72
WMT14			
En-De	27.41	27.82	28.01
De-En	32.69	33.65	34.19
ASPEC			
En-Ja	54.51	55.46	55.00
Ja-En	30.77	31.23	31.29

BPE-Dropout: Simple and Effective Subword Regularization

#token / #substring is smoother



BPE-Dropout: Simple and Effective Subword Regularization

Embedding is more reasonable

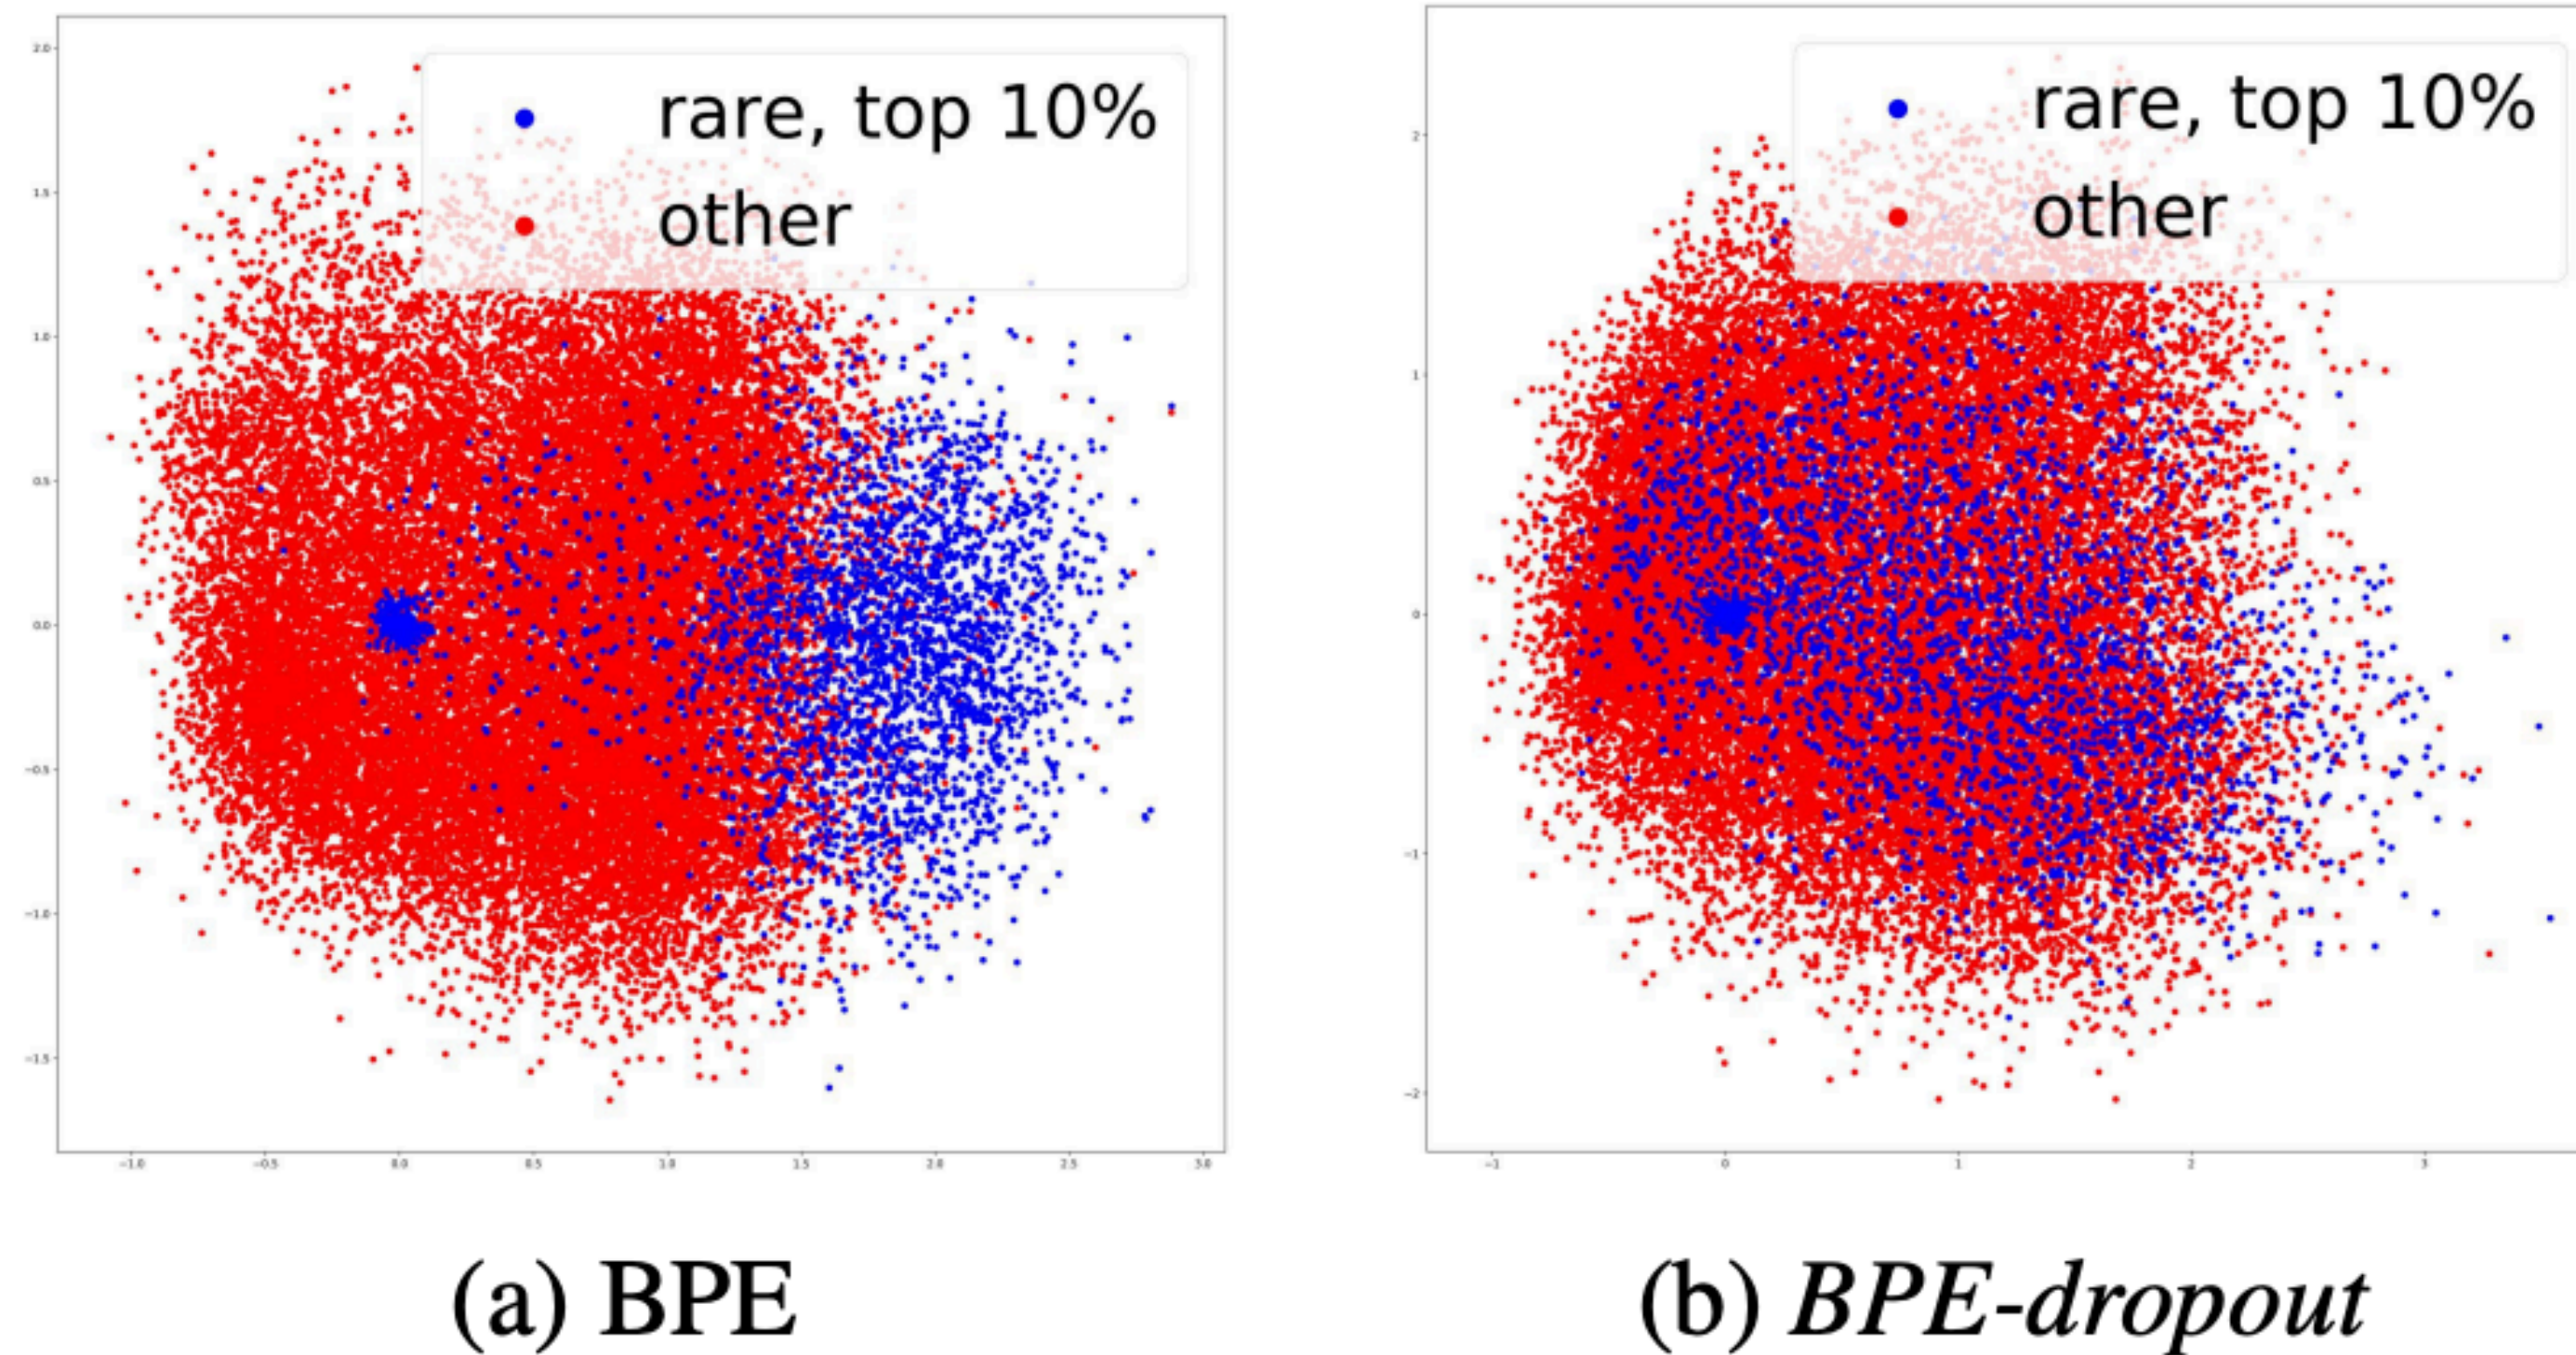


Figure 7: Visualization of source embeddings. Models trained on WMT14 En-Fr (4m).

BPE-Dropout: Simple and Effective Subword Regularization

More robust to misspelled input

source	BPE	<i>BPE-dropout</i>	diff
En-De			
original	27.41	28.01	+0.6
misspelled	24.45	26.03	+1.58
De-En			
original	32.69	34.19	+1.5
misspelled	29.71	32.03	+2.32
En-Fr (4m)			
original	33.38	33.85	+0.47
misspelled	30.30	32.13	+1.83
En-Fr (16m)			
original	34.37	34.82	+0.45
misspelled	31.23	32.94	+1.71

Conclusion

	Easy to Implement	Contribution
Subword Regularization	No	Induce probability to subword segmentation
Adversarial Subword Regularization	No	Integrate adversarial into subword segmentation
BPE-dropout	Yes	Introduce randomness in BPE merge operations

Thank You